

Introduction to My Current and Future Work

Krzysztof Drewniak

Visiting student at HPAC

4 Oct 2017

Outline

About Me

`gemm3()`

Simplifying Expressions in Linnea

About Me

Why I'm Here

- ▶ Bachelor's student in Computer Science and (Pure) Mathematics at the University of Texas at Austin
- ▶ Degree almost complete — need to write and defend research thesis to graduate with honors
- ▶ Invited by Prof. Bientinesi to work on Linnea project
- ▶ Will be staying at RWTH from September to mid-January

Research Background

- ▶ Computational epidemiology with Dr. Armin R. Mikler at U. of North Texas while in early college high school
 - ▶ Worked on how to reduce epidemic spread
 - ▶ Vaccination strategies based on centrality (in the graph that served as a model)
- ▶ Joined Dr. Robert Van De Giejn's group at UT Austin last year
 - ▶ Was initially interested in FLAME
 - ▶ Created an algorithm for $D+ = ABC$
- ▶ Interests are in the intersection of program correctness and low-level systems programming

Industry Background

Last four summers were spent doing internships:

- ▶ Google Summer of Code — adding more thorough Unicode support to a Lisp implementation
- ▶ WhatsApp — Server-side media handling improvements and some encryption things
- ▶ TrueCar — Machine learning (that is, database scouring) for fake-order detection
- ▶ Microsoft, Bing Maps — Deep learning to improve maps suggestions

gemm3()

gemm3()

Problem

- ▶ Want to compute $D+ = ABC$
- ▶ If we only have `gemm()` (or `trmm()` ...), need to break into

$$T = BC$$

$$D+ = AT$$

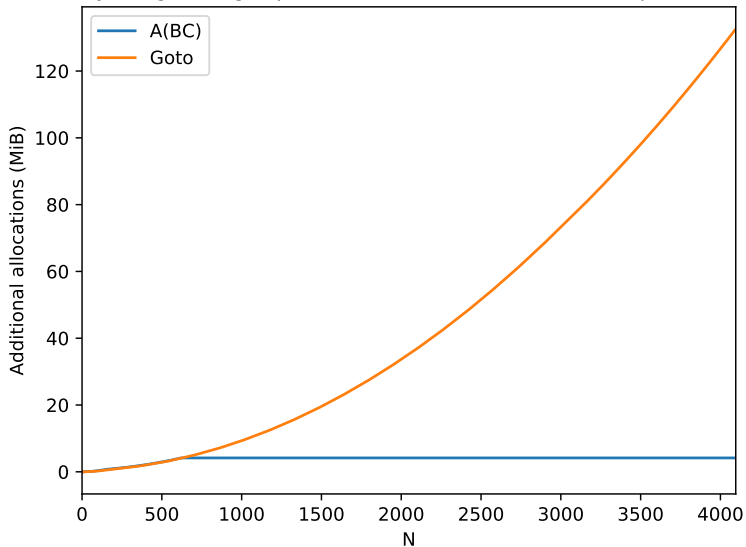
- ▶ T is often large, causing memory consumption issues
- ▶ Writing to/from memory can also be costly

Solution

- ▶ New algorithm to compute $D+ = A(BC)$ in constant memory
- ▶ Uses two different matrix multiply algorithms
- ▶ Inner algorithm is called to compute at most cache-sized blocks of (BC)
- ▶ Reduces memory impact by never storing all of (BC)
- ▶ Also increases performance by keeping intermediates in cache

Memory results


Memory usage of right parenthesized kernel vs. Goto, square matrices



Performance results

4.7% GFlops/s higher for square matrices (5.1% in left-parenthesized case)


$D += A(BC)$, square matrices



Simplifying Expressions in Linnea

Overview of Linnea

- ▶ Compiles linear algebra expressions into series of BLAS/LAPACK calls
- ▶ Generates code which significantly outperforms other systems
- ▶ Operates on symbolic expressions



- ▶ Simplifying expressions is central to Linnea
Ex. $ABB^{-1} \rightarrow AI \rightarrow A$

Simplifying expressions in Linnea

- ▶ Linnea needs to simplify expressions in order to check their equivalence and correctly apply kernels
- ▶ The simplification rules are based on linear algebra identities
 - ▶ $\alpha A + \beta A \rightarrow (\alpha + \beta)A$
 - ▶ $A^T \rightarrow A$ (only when A is symmetric)
- ▶ Simplification code is currently handwritten, requiring much time and effort

Goal

- ▶ Find a way to autogenerate the simplification module for Linnea
- ▶ Benefits:
 - ▶ Requires only knowledge of the underlying identities
 - ▶ Less hand-written code, so fewer bugs
 - ▶ Generalizability — same code could be used for a tensor compiler or similar
 - ▶ Potentially improved performance

Generating term rewriting systems

- ▶ Term rewriting systems are a set of rewrite rules that operate on algebraic expressions
- ▶ Pattern matching identifies applicable rules
- ▶ Algorithms exist to turn sets of identities into rewrite systems that produce a normal form
 - ▶ Main example: Knuth-Bendix algorithm
- ▶ Do not always succeed or even terminate
- ▶ Not known to handle constraints (like $A = A^T$ if A is symmetric)

Plan

- ▶ Adapt Knuth-Bendix (or similar) algorithms for this purpose and implement them
 - ▶ If this fails, investigate other potential techniques for this problem
- ▶ Translate known identities into rewrite system
- ▶ Generate efficient simplification code based on rewrite rules